

Chapter 10

C++ Dynamic Memory

Objectives

- To understand the similarities and differences between C++ pointers and Python references.
- To learn how to use the C++ operators that access memory addresses and dereference pointers.
- To understand how to dynamically allocate and deallocate memory in C++.
- To learn how to write classes in C++ that allocate and deallocate dynamic memory.

10.1 Introduction

As we briefly discussed in earlier chapters, the internal mechanisms that Python and C++ use for storing data in variables and names are different. In this chapter we will discuss these differences in detail. C++'s default mechanism for storing variables is different than Python's, but C++ does support pointer variables that are similar to Python references. C++ programmers can choose which mechanism to use depending on the efficiency and capabilities they need. C++ pointers give us the flexibility to delay memory allocation decisions until run-time. This makes it possible to change the size of arrays at run-time and create linked structures in C++. Using C++ pointers does require much more care than using Python references; it is easy to make mistakes with pointers and create a program that gives unexpected

results or crashes. This chapter and the next chapter will cover the use of dynamic memory and pointers. We will begin by reviewing the basic memory models of Python and C++.

Python names are a reference to a memory location where the actual data is stored along with type information and a reference count; different names can refer to the same data object and assignment statements make the name refer to a different data object. C++ associates (binds) a memory location with each variable and the same memory location is used for that variable throughout the lifetime of the variable. Each assignment statement causes different data to be stored in the memory location bound to the variable. Here is a C++ example:

```
// memory.cpp
#include <iostream>
using namespace std;

int main()
{
    int x, y, z;
    x = 3;
    y = 4;
    z = x;
    x = y;
    cout << x << " " << y << " " << z << endl;
    return 0;
}
```

The following table shows a representation of memory while this program is executing. When the `main` function begins execution, four bytes are allocated for each of the three integers. We have started our table at the memory location 1000, but the specific memory address used is not important and can vary each time the program is run. The key point to notice is that the memory location used for each variable does not change; the data stored at the memory location does change as different values are assigned to the variable. As you would expect, the program outputs 4 4 3.

Memory address	Variable name	Data value
1000	x	3 then 4
1004	y	4
1008	z	3

The Python version of this program is the following:

```
# memory.py
x = 3
y = 4
z = x
x = y
print x, y, z
```




Figure 10.1: Picture of Python memory references

The end result of executing comparable code for C++ variables and Python references to immutable types is the same even though the internal representations are different. The Python program also outputs 4 4 3. Figure 10.1 shows a pictorial representation of the memory for this Python code. The key point to notice here is that there is one copy of the 3 object and one copy of the 4 object at fixed memory locations; the names refer to these objects and as the code executes, the memory location that x refers to changes from 1000 to 1012. At the end, we have multiple names referring to the same memory location. We have not shown the reference count and type information for the object, but each Python integer object requires 12 bytes on 32-bit systems.

The important differences between Python and C++ are

- Each C++ variable corresponds to a fixed memory location where data is stored; each time a value is assigned to that variable, the same memory location is used to store the data.
- A Python name refers to an object in memory. Python objects must also store information about their type and a reference count, so storing data in Python requires more space than storing the same data in C++.

In Python, both `r1` and `r2` refer to the same object that we will assume is stored at memory location 1000, as the following table shows. Additional memory is required in Python to store information about the data type for the object and the reference count for the object, but we will not include that here. Since we are creating a `Rational` object that has the instance variables `num` and `den`, we have named the variable name column based on the instance variable names of the `Rational` object.

Memory address	Variable name	Data value
1000	.num	reference to 2, then 1
1004	.den	reference to 3

If you recall our discussions from section 4.2, you will understand that the sample Python code will change the one `Rational` object to which both names refer.

The corresponding C++ declaration of `r1` and `r2` results in two `Rational` objects being created, requiring a total of 16 bytes of memory being allocated as the following table shows. Each `Rational` object requires eight bytes since it has two `int` instance variables that each require four bytes. No additional memory is needed in C++ since the C++ run-time environment does not need to keep track of the data type or the reference count.

Memory address	Variable name	Data value
1000	r1.num	?
1004	r1.den	?
1008	r2.num	?
1012	r2.den	?

After the statement `r1.set(2, 3)`, the memory now holds

Memory address	Variable name	Data value
1000	r1.num	2
1004	r1.den	3
1008	r2.num	?
1012	r2.den	?

Unless we have defined our own `operator=` (we will discuss this in subsection 10.4.3) for the C++ `Rational` class, the execution of `r2 = r1` effectively causes the two statements `r2.num = r1.num` and `r2.den = r1.den` to be executed. We cannot explicitly write those two statements ourselves since the instance variables are

Python's memory management mechanism is known as *implicit heap dynamic*. The Python interpreter automatically allocates and deallocates memory as needed. A section of memory known as a *dynamic memory heap* (sometimes referred to simply as a *heap*) is used for these allocations and deallocations. The default C++ memory management mechanism is known as *stack dynamic*. When a function is called, the amount of space needed for the variables is allocated on a stack. Since in most cases we can determine at compile time how much memory is needed for all the local variables, one machine language instruction can be used to allocate the space on the stack. When the function ends, the stack shrinks back to the space it was before the function call, effectively deallocating the memory for the local variables.

As you might have figured out by now, C++ must support another mechanism for allocating memory for variables since the Python interpreter is written in C. C++'s other technique is known as *explicit heap dynamic*. Like Python, a section of memory known as the *dynamic memory heap* (or just *heap*) is used for these allocations and deallocations. However, as the term explicit heap dynamic implies, in C++ your program code must include instructions that directly allocate and deallocate the memory. C++ uses pointer variables to support this dynamic memory allocation and deallocation. With C++ pointers we can write code that allows us to determine and change the amount of memory allocated at run-time (rather than setting the amount at compile time). We can write data structures that can grow in size as needed and write linked structures using C++ pointers. This chapter will discuss how to use C++ pointers, how they are similar to Python references, and how to write C++ classes that use dynamic memory. We will learn how to write linked structures in C++ in Chapter 11.


```
int main()
{
    int *b, *c, x, y;
    x = 3;
    y = 5;
    b = &x;
    c = &y;
    *b = 4;
    *c = *b + *c;
    cout << x << " " << y << " " << *b << " " << *c << " ";
    c = b;
    *c = 2;
    cout << x << " " << y << " " << *b << " " << *c << endl;
    return 0;
}
```

The unary ampersand operator computes the address of its operand. Thus, the statement `b = &x` causes the program to store the memory address of `x` in the memory for the variable `b`. The following table indicates that the computer used memory addresses 1000 through 1015 to store our variables and shows the value of each variable after the statement `c = &y` is executed. The computer does not necessarily use the address starting at 1000, but we commonly use that address in our examples in this book.

Memory address	Variable name	Data value
1000	b	1008
1004	c	1012
1008	x	3
1012	y	5

The unary asterisk operator is used to *dereference* a pointer. Dereferencing a pointer means to access the data at the address the pointer holds. The statement `*b = 4` causes the program to store the data value 4 at memory address 1008 (since `b` currently holds 1008). Based on this knowledge, see if you can determine the output of the sample program before reading the next paragraph.

The statement `*c = *b + *c` determines the integer values at memory address 1008 (the address `b` points to) and memory address 1012 (the address `c` points to) and adds the 4 and 5 together. The result, 9, is stored at memory address 1012 (the address that `c` points to). The statement `c = b` copies the data value for `b`, which is the address 1008, to the memory for `c` (i.e., 1008 is now stored at memory location 1004). You should note that assigning pointer variables is essentially the same as assigning two names in Python; both `b` and `c` now refer to the same data. Based on

The C++ code using C++ pointers that corresponds to the same Python `Rational` example earlier in the chapter is the following

This example outputs 1/3 for **r1** and 1/3 for **r2** since **r1** and **r2** are pointers to the same memory locations. The memory table for this code fragment is:

The declarations of `r1` and `r2` result in four bytes being allocated for each one since pointers require four bytes. The `Rational` constructor is not called when you declare a pointer since we are creating a pointer, not a `Rational` object. The statement `r1 = new Rational` results in eight bytes being allocated since the two integer instance variables `num_` and `den_` require a total of eight bytes. The `r1 = new Rational` statement also causes 2000 to be stored in the memory location for variable `r1`. The constructor is called by `r1 = new Rational` since it creates a `Rational` object. The `r1->set(2, 3)` statement results in 2 being stored at memory location 2000 and 3 being stored at memory location 2004.


```
d = new double[n];
for (i=0; i<n; ++i) {
    cout << "Enter number " << i << ": ";
    cin >> d[i];
}
delete [] d;
```

The example allows the user to specify the array size at run-time. The **new** command allocates the specified amount of memory and returns the starting address of the allocated memory. When the brackets (`[]`) are used after the data type in the **new** statement, the amount of memory necessary to store the number of items specified inside the brackets is allocated and the starting address is returned. In this case, `n*8` consecutive bytes would be allocated on machines that use eight bytes to store a double value. The expression inside the brackets indicates how many values of the type `double` to allocate; an array of size `n` was allocated so the valid index values are 0 through `n-1`. After the dynamic memory has been allocated, it can be accessed using the array bracket notation. The same index array calculations discussed in section 8.11 can be used since the pointer variable holds the starting address of a contiguous section of memory.

Whenever you allocate memory dynamically, you must also deallocate the memory with a statement that executes later in your program. Since we allocated an array, we must tell the `delete` statement to deallocate an array instead of the memory that holds a single value. The square brackets are used with both the `new` statement and the `delete` statement when allocating and deallocating arrays. You do not indicate the size of the array when deallocating a dynamic array; the C++ run-time environment knows how much memory to deallocate. Repeatedly allocating memory and forgetting to deallocate memory in a C++ program will eventually result in your program using up a large percentage of the computer's memory, causing the computer to slow as it uses the hard disk for extra memory. This is why it is important to deallocate memory when it is no longer needed.

The main reason for using dynamic arrays is that you do not need to know the size of the array at compile time. In many cases, you still may not know the size needed when the array is first allocated. The Python built-in list allows you to append as many items as you want so there is no need to determine how much memory to allocate the first time you allocate memory; it would be impossible to anticipate how much memory to allocate ahead of time since different uses of the list will require different sizes. Once the array fills up, we may need to make the array larger. Because the memory immediately following the dynamic array may already be in use (remember that array elements must be in consecutive memory locations), we cannot make the array larger. The solution is to allocate a new larger

array, copy the values from the original array to the new array, and then delete the original array. The following code fragment demonstrates this:

```
int *data, *temp;
int i;

// create original array
data = new int[5];
for (i=0; i<5; ++i) {
    data[i] = i;
}
// create new larger array
temp = new int[10];
// copy from original array to larger array
for (i=0; i<5; ++i) {
    temp[i] = data[i];
}
// deallocate original array
delete [] data;
// make data point to new larger array
data = temp;
// now we can access positions 0-9
for (i=5; i<10; ++i) {
    data[i] = i;
}
// deallocate last allocation
delete [] data;
```

The memory table for this code after the first **new** statement and **for** loop are executed is below. We will assume the memory addresses used for the local variables start at memory location 1000 and that the dynamically allocated memory is the block of memory from 2000 through 2019 (four bytes for each of the five integers).

Memory address	Variable name	Data value
1000	data	? then 2000
1004	temp	?
1008	i	5
2000		0
2004		1
2008		2
2012		3
2016		4

After the memory is allocated for the `temp` pointer, the values are copied from the original array, and the values 5 through 9 are stored in the larger array, the memory table is the following assuming the memory starting at location 3000 is used for the `temp` pointer.

Memory address	Variable name	Data value
1000	data	2000
1004	temp	3000
1008	i	10
2000		0
2004		1
2008		2
2012		3
2016		4
3000		0
3004		1
3008		2
3012		3
3016		4
3020		5
3024		6
3028		7
3032		8
3036		9

After the first `delete [] data` statement, the memory at locations 2000–2019 is deallocated and returned to the dynamic memory heap so it can be used again. The statement `data = temp` stores 3000 at memory location 1000 (i.e., `data` now points to the second larger allocated array). At this point, both `data` and `temp` point to the same dynamically allocated array. This is the same concept as having two references to the same data in Python. After that assignment statement, the next loop fills in the values 5 through 9 in memory locations 3020 through 3039. The final `delete [] data` statement then deallocates the memory locations 3000–3039 so they can be used again.

Figure 10.2 shows a pictorial representation of this. The top part of the figure shows the representation after we have created the new larger array and copied the values from the first array. The middle part of the figure shows the state after the

first `delete data` statement. The bottom part of the figure shows the state just before the final `delete [] data` statement.




Figure 10.2: Pictorial representation of resizing a dynamic array

If we then fill up this new array and need a larger array, we allocate a larger array, copy the values from the previously allocated array, and then delete the previous array. Each resizing operation results in the previous array being deleted so that we do not have a memory leak if we perform this resizing operation multiple times. In our example, `data` points to the last array that was allocated (once we execute the `data = temp` statement). This pattern of allocating a new section of dynamic memory using a different pointer variable, copying the values from the old section to the new section, deallocating the old section, and setting the original pointer variable to the new section is a common pattern in C++ dynamic memory, so make certain you fully understand how it works and why the order of the steps is important. In the next section, we will examine this pattern using a class.

10.4 Dynamic Memory Classes

When you write a class that dynamically allocates memory for pointer instance variables, you need to make certain that the memory is properly deallocated. There are three additional C++ methods that dynamic memory classes must use to properly allocate and deallocate memory. These three methods are the destructor, copy constructor, and assignment operator (`operator=`). If your class does not use dynamic memory, you do not need to write any of these methods. Classes that use dynamic memory must write a destructor that deallocates the memory. The other two methods may be implemented or declared in the `private` section but not implemented. Declaring them in the `private` section but not implementing them prevents them from being called. We will discuss the details of when these methods are called and what they must do in this section. Implementing them correctly will prevent your class from having memory leaks or other memory errors.

10.4.1 Destructor

As discussed in the previous sections, in C++ you must explicitly deallocate any memory that you explicitly allocate with the `new` command. C++ classes have a special method known as a *destructor* that is used for deallocating memory. The destructor method has the same name as the class, with a tilde (~) in front of it. Just as constructors do not have a return type, the destructor also does not have a return type. The purpose of the destructor is to deallocate any dynamic memory the class has allocated that has not yet been deallocated. You never directly call the destructor using the name of the method; it is called automatically when an instance of the class goes out of scope or when you use the `delete` operator on a pointer to an instance of the class. If your class uses dynamic memory and does not have a destructor, your code will in most cases have a memory leak.

We will start with a simple dynamic array class that we will extend throughout this chapter to demonstrate how to correctly write dynamic memory classes. In this first version of the class, we will write all the methods inline in the header file. We have added a few output statements so you can see that the constructor and destructor are called. This `List` class uses three instance variables. The instance variable `data_` is used to hold the starting address of the dynamic array containing the list's values. The `size_` instance variable indicates how many items are currently in the list. The `capacity_` instance variable indicates how large the dynamic array is (i.e., how many items the list can hold before the dynamic array needs to be resized).

```

// List1.h
#ifndef __LIST_H__
#define __LIST_H__

#include <iostream>

class List {
public:
    List(int capacity=10);
    ~List() { delete [] data_; std::cout << "destructor\n";}

private:
    int size_;
    int capacity_;
    int *data_;
};

inline List::List(int capacity)
{
    std::cout << "constructor\n";
    data_ = new int[capacity];
    size_ = 0;
    capacity_ = capacity;
}

#endif __LIST_H__

```

We did not put the `using namespace std` statement in the header file since any file that included the header file would have this statement. Also note that we put the default value for the parameter `size` only in the declaration of the `List` constructor and not in the implementation of the constructor. Here is a simple program that uses our class:

```

// test_List1.cpp
#include <iostream>
using namespace std;

#include "List1.h"

int main()
{
    List b;
    return 0;
}

```

When this program is compiled and executed, it outputs

```

constructor
destructor

```

The declaration `List b` causes the constructor to be called and it allocates dynamic memory. At the end of the `main` function, the variable `b` goes out of scope so the destructor method is automatically called and it deallocates the dynamic memory. This is why the program outputs the two lines. The following program also causes the same output.

```
// test_Listp.cpp
#include <iostream>
using namespace std;

#include "List1.h"

int main()
{
    List *b; // constructor is not called here

    b = new List(20); // constructor is called here
    delete b; // destructor is called here
}
```

The comments indicate when the constructor and destructor are called. Remember the declaration `List *b` causes four bytes that can store an address to be allocated. The `new` statement causes a `List` object to be created by calling the constructor with the specified size. The `delete` statement causes the `List`'s destructor to be called and the dynamic memory the constructor allocated is deallocated. When the variable `b` goes out of scope at the end of the `main` function, the four bytes for the pointer are automatically deallocated just as the memory for any variables are when they go out of scope. This is the reason you only need to write a destructor when your class allocates dynamic memory.

10.4.2 Copy Constructor

As the name implies, the purpose of a *copy constructor* is to create a new object by copying an existing object. In C++, the copy constructor for a class is called when you pass an instance of a class by value to a function or method. Remember that pass by value requires that a separate copy of the actual parameter be created. You can also call a copy constructor directly when declaring a variable as we will demonstrate later in this section.

Unless you write a copy constructor for a class, the C++ compiler generates a default copy constructor for you. The default copy constructor it creates effectively

we only copied up to the value of `size_` since the values past that are not relevant to the object. At this point we do not have any way of putting elements in our simplified class to cause `size_` to be more than zero, but our final version will. The following example using this class allows us to see when the copy constructor is called.

```
// test_List2.cpp
#include <iostream>
using namespace std;
#include "List2.h"

void f(List c)
{
    cout << "start f\n";
    cout << "end f\n";
}

void g(List &d)
{
    cout << "start g\n";
    cout << "end g\n";
}

int main()
{
    List b;
    f(b);
    g(b);

    List e(b);
    return 0;
}
```

The output of the program is the following. The comments in parentheses are obviously not part of the output, but explain what caused each of the methods to be called.

```
constructor (create b in main function)
copy constructor (create the copy c from b in f function)
start f
end f
destructor (destructor for c when function f completes)
start g
end g
copy constructor (create e in main function)
destructor (destructor for e or b)
destructor (destructor for e or b)
```


As we mentioned earlier, you can declare the copy constructor private and not implement it. This prevents code that uses your class from causing the copy constructor to be called; the code would not be able to pass an instance of your class by value to a function or method or explicitly call the copy constructor. Code that attempts to perform either of those actions will generate a compiler error. If your class uses a large amount of memory, you may want to do this to prevent a user of your class from making a copy of it. The following header file demonstrates this.

```
// List3.h
#ifndef __LIST_H__
#define __LIST_H__

#include <iostream>

class List {
public:
    List(int capacity=10);
    ~List() { delete [] data_; std::cout << "destructor\n"; }
private:
    List(const List &source);
    int size_;
    int capacity_;
    int *data_;
};
```



```
private:
    int size_;
    int capacity_;
    int *data_;
};

inline List::List(int capacity)
{
    data_ = new int[capacity];
    size_ = 0;
    capacity_ = capacity;
}

inline List::List(const List &source)
{
    int i;

    size_ = source.size_;
    capacity_ = source.capacity_;
    data_ = new int[capacity_];
    for (i=0; i<size_; ++i) {
        data_[i] = source.data_[i];
    }
}

inline void List::operator=(const List &source)
{
    int i;

    if (this != &source) {
        delete [] data_;
        size_ = source.size_;
        capacity_ = source.capacity_;
        data_ = new int[capacity_];
        for (i=0; i<size_; ++i) {
            data_[i] = source.data_[i];
        }
    }
}

#endif __LIST_H__
```

Since the object has already been created, the assignment operator is a little more complicated than the copy constructor. We must properly deallocate memory that has already been allocated and ensure that the class is not accidentally assigning the object to itself or we will deallocate the only copy of the data. In C++ classes, the identifier `this` is an implicit pointer to the object with which the method is

The bracket operator (`operator[]`) provides the same functionality as the Python `__getitem__` and `__setitem__` methods. It is declared inline after the class definition and demonstrates a reference return type. The ampersand after the type name indicates that a reference to an integer is returned, meaning it effectively returns the address of the position in the array. This allows the operator to be used on the left-hand side of an assignment statement as `b[0] = 5` where `b` is an instance of our `List` class. It can also be used on the right-hand side of an assignment statement or as part of an expression just as a non-reference return type can. Without the reference return type, the operator could only be used on the right-hand side of an assignment statement (corresponding to only the Python `__getitem__` method). Returning a reference only makes sense if it is a reference to an instance variable or dynamically allocated memory. We will discuss this later in the chapter.

The `List` class provides an array of integers whose initial size is specified when the constructor is called. The constructor allocates a dynamic array with the specified capacity and initializes the `size_` instance variable to indicate the list is empty. If we did not allocate the memory in the constructor, but instead deferred it to another method (such as the first time the `append` method is called), we would initialize the pointer variables to `NULL`. The `NULL` constant is defined in the `cstdlib` header file. The value `NULL` is defined to be zero which is never a valid address for memory that has been dynamically allocated. The use of `NULL` in C++ to indicate an invalid pointer is similar to the use of `None` in Python to indicate a reference that is not initialized to an object of a specific type.

The class makes use of a private method named `copy` to implement the code that is needed in both the copy constructor and assignment operator. The assignment operator needs extra code to deallocate the existing dynamic array before allocating a new dynamic array of the appropriate size and copying the data. Remember that a copy constructor is creating a new object, so no memory has been previously allocated for the object when the copy constructor is called. However, the variable on the left-hand side of an assignment statement has already had its constructor called and memory allocated, so that needs to be deallocated. The code for the destructor follows the copy constructor. The destructor simply deallocates the dynamic array and is called automatically when a non-pointer instance of `List` goes out of scope. Note that it does not deallocate the non-pointer instance variables since the memory for those is automatically deallocated. In this example, we are using a separate implementation file unlike the earlier simplified examples in which the entire class was written in the header file.

```
// List.cpp
#include "List.h"

List::List(size_t capacity)
{
    data_ = new int[capacity];
    capacity_ = capacity;
    size_ = 0;
}

List::List(const List &list)
{
    copy(list);
}

List::~List()
{
    delete [] data_;
}
```

We have written the `operator=` method slightly differently so that we can use it in a chained assignment statement. The method returns a reference to a `List` object. By returning `*this`, we are returning the `List` object that we just assigned. This allows us to write the chained form of the assignment statement (e.g., `b = c = d`). Remember that the assignment operator is right to left so it is equivalent to `c = d; b = c`. By returning a reference to the left-hand parameter, the result of `c = d` is the object `c` that we then use as the right-hand parameter when assigning `b`.

The `copy` method that is used by both the `operator=` and the copy constructor allocates an array of the same size as the `List` object that is passed to it and copies all the data from the parameter object's array into the newly allocated array. We have also added the `operator+=` method so we can demonstrate another potential pitfall.

```
void List::copy(const List &list)
{
    size_t i;
    size_ = list.size_;
    capacity_ = list.capacity_;
    data_ = new int[list.capacity_];
    for (i=0; i<list.capacity_; ++i) {
        data_[i] = list.data_[i];
    }
}
```


of allocating a new larger array, copying the data to it, updating the pointer, and then deallocating the old smaller array, as we discussed in section 10.3.

```
void List::append(int item)
{
    if (size_ == capacity_) {
        resize(2 * capacity_);
    }
    data_[size_++] = item;
}

// should this method have a precondition? see end of chapter exercises
void List::resize(size_t new_size)
{
    int *temp;
    size_t i;

    capacity_ = new_size;
    temp = new int[capacity_];
    for (i=0; i<size_; ++i) {
        temp[i] = data_[i];
    }
    delete [] data_;
    data_ = temp;
}
```

We leave it as an exercise to add the other methods in the built-in Python list's API to this C++ dynamic memory list. As we mentioned earlier, you do need to be careful when using `unsigned int` or the equivalent `size_t` data type. As we listed in Figure 8.4, the range of the `int` type on 32-bit systems is usually from about negative two billion to about positive two billion while the `unsigned int` type ranges from zero to about four billion. With the `unsigned int` data type, there is no bit representation that corresponds to a negative number. So the question is, what happens when an operation would result in a negative number?

```
// unsigned.cpp
#include <iostream>
using namespace std;
int main()
{
    unsigned int x = 0;
    x--;
    cout << x << endl;
    return 0;
}
```


pages your program is using, a hardware exception is generated and the program crashes.

Since the hardware detects errors only when a program attempts to access memory that is not one of the pages the operating system allocated for the program, a program may access memory that is one of its pages but is not a valid memory address that it should be using. In this case, your program will not crash immediately, but it can have different results each time it runs or it can crash at a later point in time.

We will start with a simple example that does not use dynamic memory, but could give unexpected results. See if you can find the error in this program before reading the paragraph after the code.

```
// this program is incorrect
#include <iostream>
using namespace std;

int main()
{
    int x[10];
    int y = 0;
    int i;

    for (i=0; i<=10; ++i) {
        x[i] = i;
    }
    cout << "y=" << y << endl;
    return 0;
}
```

Unlike Python, C++ does not do any index checking when you attempt to access an element in an array (Python does check when you attempt to access an element in a list). The problem with this example is the array can be indexed using the values 0 through 9, but the `for` loop sets `x[10]`. Depending on how the memory is allocated for the local variables, this could result in the program outputting 10 for `y` even though we set `y` to zero. If the memory location used for the variable `y` is immediately after the memory location for the array, then `x[10]` and `y` correspond to the same memory address. If the memory for the variables is not allocated in this order then the program produces the expected output of 0.

Remember that pointers hold an address and dereferencing a pointer attempts to read or store data at that address. This is what can lead you to access memory you should not be using. The following simple program is almost guaranteed to crash on any computer.

We will examine one more program with errors in this chapter, but there are lots of different ways you can have these problems. This program has two errors.

```
// this program is incorrect
int main()
{
    int *x, *y;
    x = new int;
    y = new int;
    *x = 3;
    y = x;
    *y = 3;
    delete y;
    delete x;
}
```

The first problem is that this program has a memory leak. The memory for two integers is allocated, but then the statement `y = x` causes both pointers to refer to the same memory location. This results in the memory allocated by the statement `y = new int` being leaked since there is no way to access it and delete it. The `delete y` statement deletes the memory allocated by the `x = new int` statement. Since `x` also pointed to that memory location, the statement `delete x` attempts to deallocate the same block of memory a second time. This will likely corrupt the dynamic memory heap. This can also cause your program to crash immediately, or at a later time, or never.

10.5.3 Memory Error Summary

Some C++ run-time environments do not show you the exact line where your program crashed or a stack trace showing the function or method calls that resulted in the program crashing at that line. Most IDEs (integrated development environments) will show you the execution traceback similar to Python indicating at what line the program crashed and the functions or methods that were called to get to that point. This information is important for determining why your program crashed. Unfortunately, as we have discussed, the line your program crashed at is not necessarily the line that is incorrect. If the line it crashes at is dereferencing a pointer, the problem is that either you forgot to give that pointer a valid address or

somehow it ended up pointing to memory that is no longer valid for your program to use (for example, you already called `delete` on that memory block or it got set to a value that does not correspond to a valid address). The traceback tells you the order the functions or methods were called to the point of the crashing. This helps you determine the code that caused the problem.

Also, as we mentioned earlier, sometimes you can corrupt the dynamic heap by accessing incorrect memory locations or calling `delete` twice for the same block of memory. This will typically not result in a crash until you try to allocate memory again. These types of errors can be extremely difficult and frustrating to track down. Fortunately, while you are developing your code, you can use an IDE that provides a debugger to help track down these errors. Debuggers provide a number of features to help you find errors in your programs. Most allow you to stop execution at specific source code lines within your program, examine the values of variables at that point, and execute one line or one function at a time while you watch the values of the variables. Debuggers typically provide additional capabilities beyond the ones we listed here.

When running your program within a debugger and your program crashes, the debugger will typically show you similar information to the Python traceback. It is fairly easy to develop Python code without a debugger, but when writing dynamic memory code in C++, a debugger and good IDE will help you track down memory errors more quickly and with less frustration. Sometimes proofreading the code around the crash (or for the entire class if the crash is in a method) is the most effective way to solve the problem.

It is always a good idea to find the smallest sample input that causes your program to crash or to work incorrectly. This is especially important when dealing with dynamic memory errors. If we determined that our `List` class did not work correctly and crashed in the `append` method, we should first check if it can happen when appending fewer items than cause the `resize` method to be called. If this is the case and we have only called the `append` method and the constructor, we know that the problem is with the constructor or `append` method. If it crashes in the `append` method only after the `resize` method has been called, then the problem could be in the constructor, `append`, or `resize`, but in this case we recommend checking the `resize` method first. Try to minimize the amount of code that is executed but still causes the problem. Limiting the amount of code you have to check will enable you to find the problem faster and with less frustration.